

Hadoop & Pig



Dr. Karina Hauser
Senior Lecturer
Management & Entrepreneurship


Outline

- Introduction
- (Setup)
- Hadoop, HDFS and MapReduce
- Pig

Introduction

- **What is Hadoop and where did it come from?**

Big Data



Big Data Sources

- **Every day 2.5 quintillion bytes or 2.5 exabytes (10^{18}) are generated, that number is estimated to double every 40 month**
- **Astronomy**
 - Sloan Digital Sky Survey (SDSS) began collecting astronomical data in 2000; 200 GB (10^9) per night
 - Large Synoptic Survey Telescope (LSST) (~2020); estimated ~20 TB (10^{12}) per night
- **Business**
 - Twitter: 12 terabytes (10^{12}) of Tweets every day
 - Walmart: 2.5 petabytes (10^{15}) of data every hour from its customer transactions

Big Data - Big Business

- **IDC predicts big data technology and services will grow worldwide from \$3.2 billion in 2010 to \$16.9 billion in 2015. This represents a compound annual growth rate of 40 percent — about seven times that of the overall information and communications technology market.**


Big Data in the News

- **The Economist Intelligence Unit study showed that nine out of 10 surveyed business leaders believe data is now the fourth factor of production, as fundamental to business as land, labor and capital.**

What is Big Data ?

- **“When the data itself becomes part of the problem”**
- **Three (to five) dimensions:**
 - Volume
 - Variety
 - Velocity
 - (Veracity)
 - (Value)

The Solution



Short History

- **Created by Doug Cutting, named after his son's toy elephant**
- **2002 - Nutch, search engine, scalability problems**
- **2004 - Google papers on GFS and MapReduce**
- **2006 - Yahoo hires Doug to improve Hadoop**
- **2008 - Hadoop becomes Apache Top Level Project**

Hadoop Today

- **Moving from “Internet” companies**
 - Yahoo
 - Google
- **to business and science applications**
 - Customer relationship management
 - Bioinformatics
 - Astrophysics

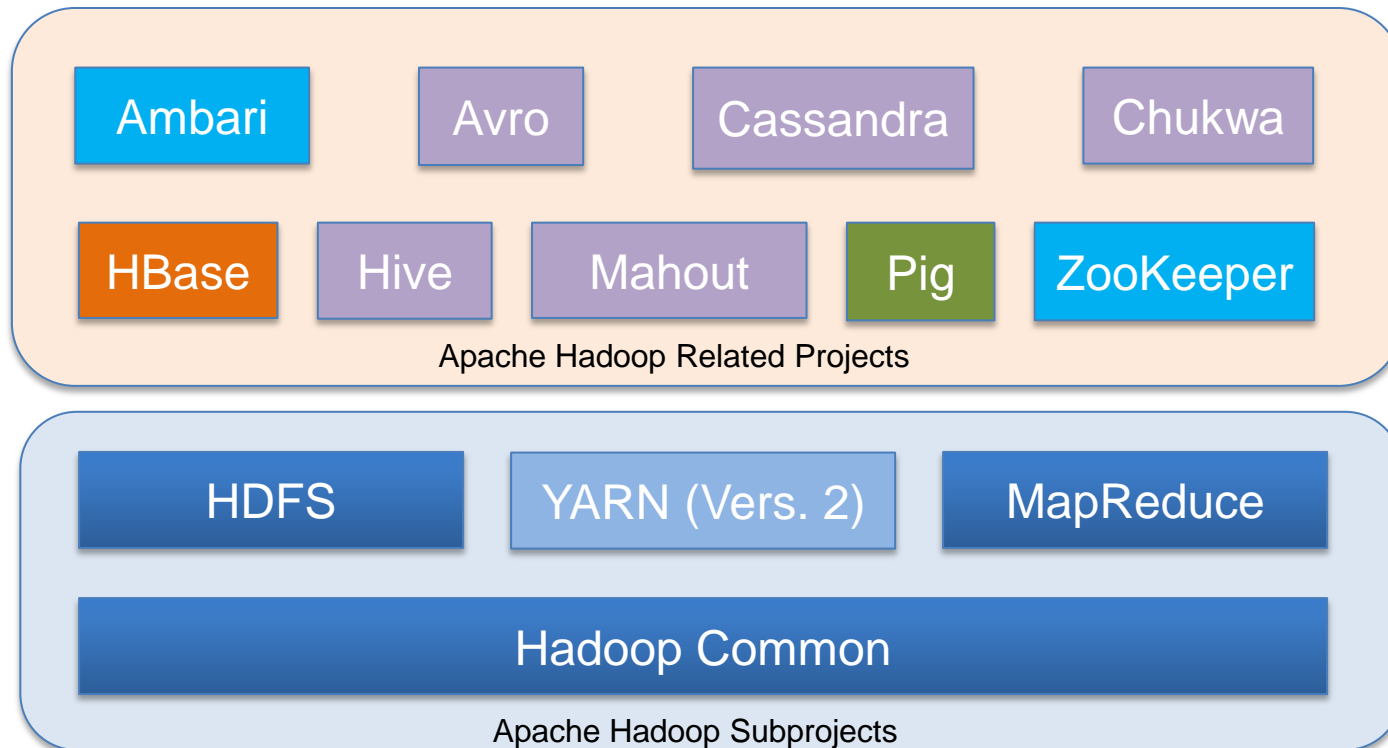


Hadoop Definition

- **“Framework that allows for the**
 - distributed processing of
 - large data sets
 - across clusters of computers
 - using a simple programming model”
- **Open-source software, maintained by**
“The Apache Software Foundation”
- **<http://hadoop.apache.org/>**



Apache Hadoop Projects




Hadoop Cluster



- **Commodity hardware**
- **Individual disk space on each node**
- **Hadoop framework handles:**
 - Data "backups" (through replication)
 - Hardware failure
 - Parallelization of code (through MapReduce paradigm)

HDFS



- **Write-once, read-many**
- **Each file is stored as sequence of same-sized blocks (default size 64MB)**
- **Blocks are replicated across different nodes**
- **Highly reliable:**
 - Redundant data storage
 - Heartbeat messages to detect connectivity problems
 - Automatic failover

HDFS

Yarn (Vers.2)

MapReduce

Hadoop Common

Apache Hadoop Subprojects

HADOOP
DISTRIBUTED
FILE
SYSTEM
(HDFS)

THE CAST

People sit in front of me
and ask me to read/write data



CLIENT

There is only ONE of me..

..and I coordinate
everything around here



NAMENODE

We store data..
..there are MANY of us
sometimes even thousands!



DATANODES




Leeds School of Business
UNIVERSITY OF COLORADO BOULDER

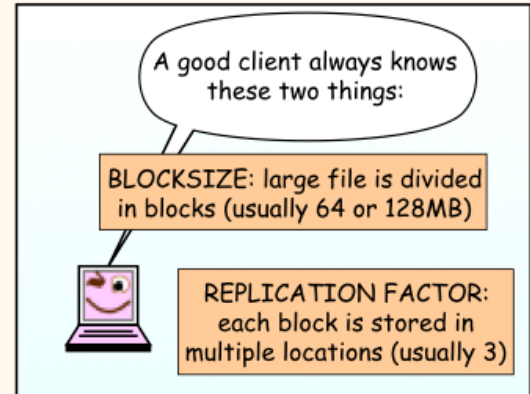

From: <http://nosql.mypopescu.com/post/15561851616/hadoop-distributed-file-system-hdfs-a-cartoon-is-worth>

WRITING DATA IN HDFS CLUSTER


REQUEST FROM USER




BLOCK AND REPLICATION



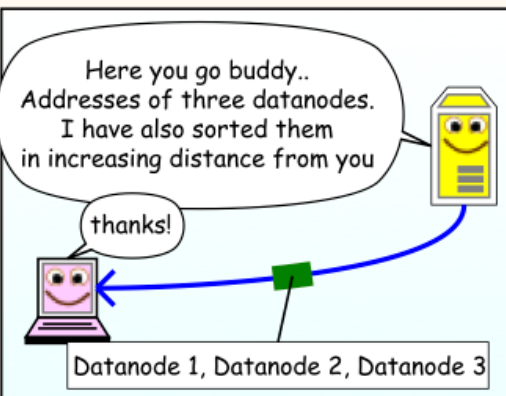

DIVIDE FILE INTO BLOCKS





ASK NAMENODE



NAMENODE ASSIGNS DATANODES



CLIENT STARTS WRITING DATA



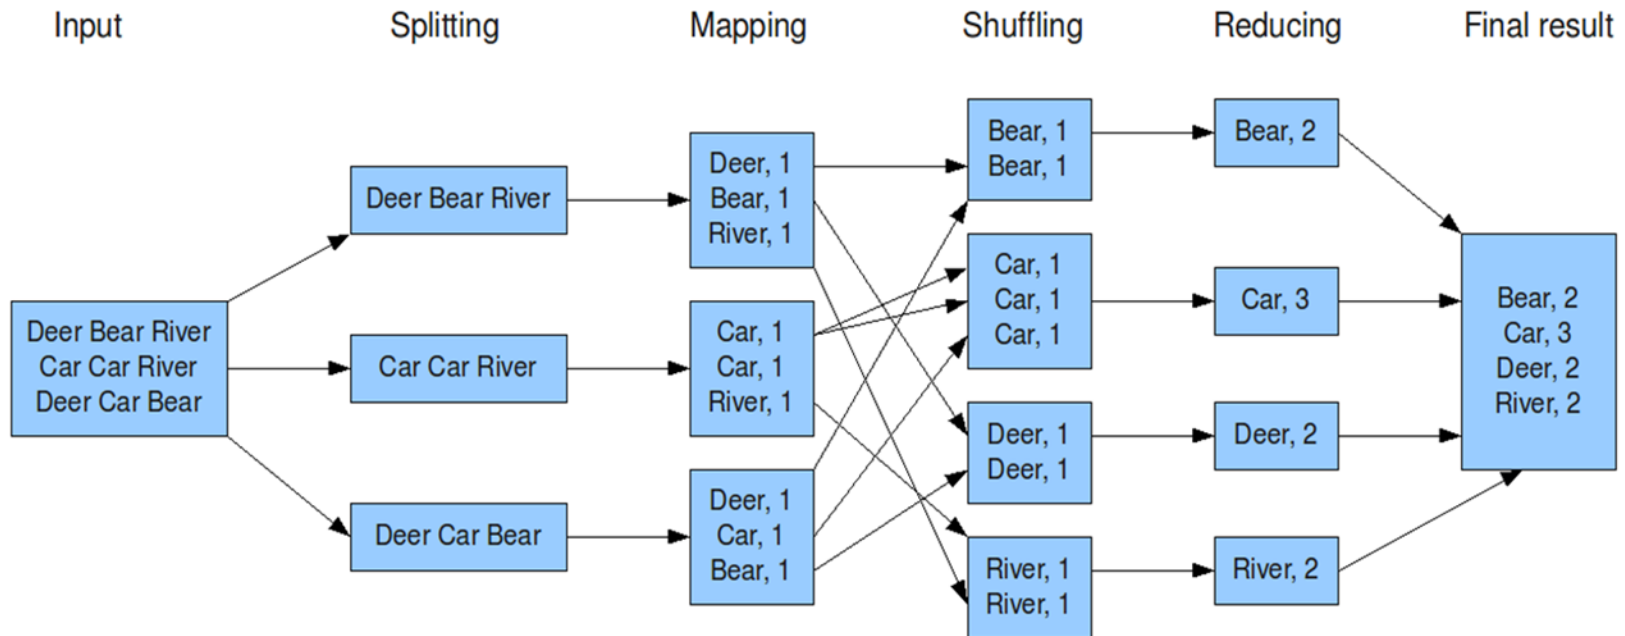
MapReduce



- **Programing model designed for**
 - batch processing of large volumes of data
 - in parallel
 - by dividing the work into a set of independent tasks
- **Not limited to Hadoop**

MapReduce WordCount Example

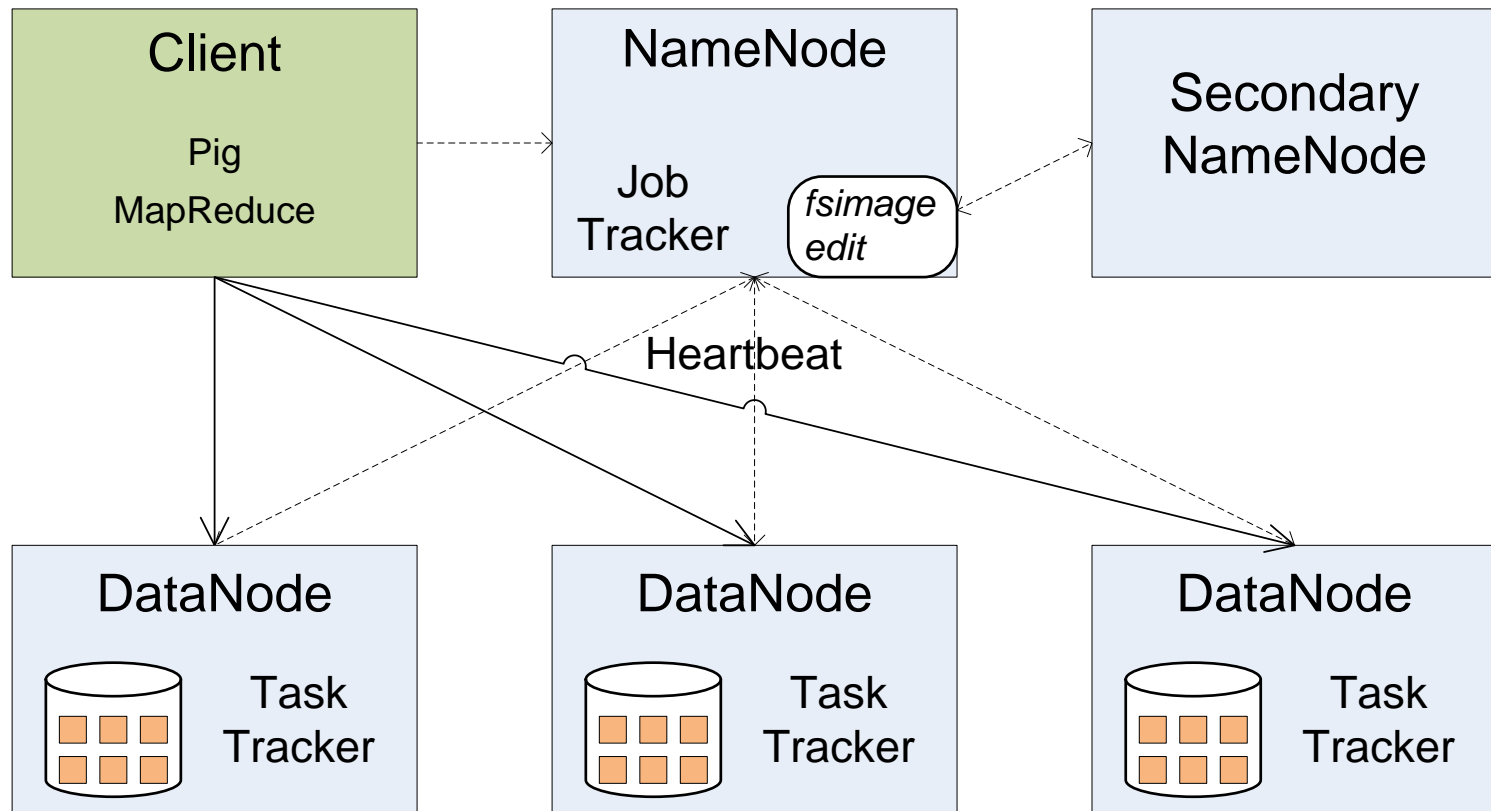
The overall MapReduce word count process



Problems suited for MapReduce

- **Iterate over a large number of records**
- **Extract something of interest from each**
- **Shuffle and sort intermediate results**
- **Aggregate intermediate results**
- **Generate final output**

Hadoop 1.x Components



Hadoop 1.x Components

- **Name Node**

- Stores metadata (filenames, replications factors ...)
 - *fsimage: latest checkpoint of namespace*
 - *edits: log of changes to namespace*
- Checks data node availability (Heartbeat)
- If possible: Separate machine

- **Data Node**

- Stores data
- Replicates blocks
- Computation

Hadoop 1.x Components

- **Secondary/Checkpoint Name Node**

- Periodically creates checkpoints of namespace
 - *Downloads fsimage and edits, creates new namespace and updates Name Node*
- Separate machine

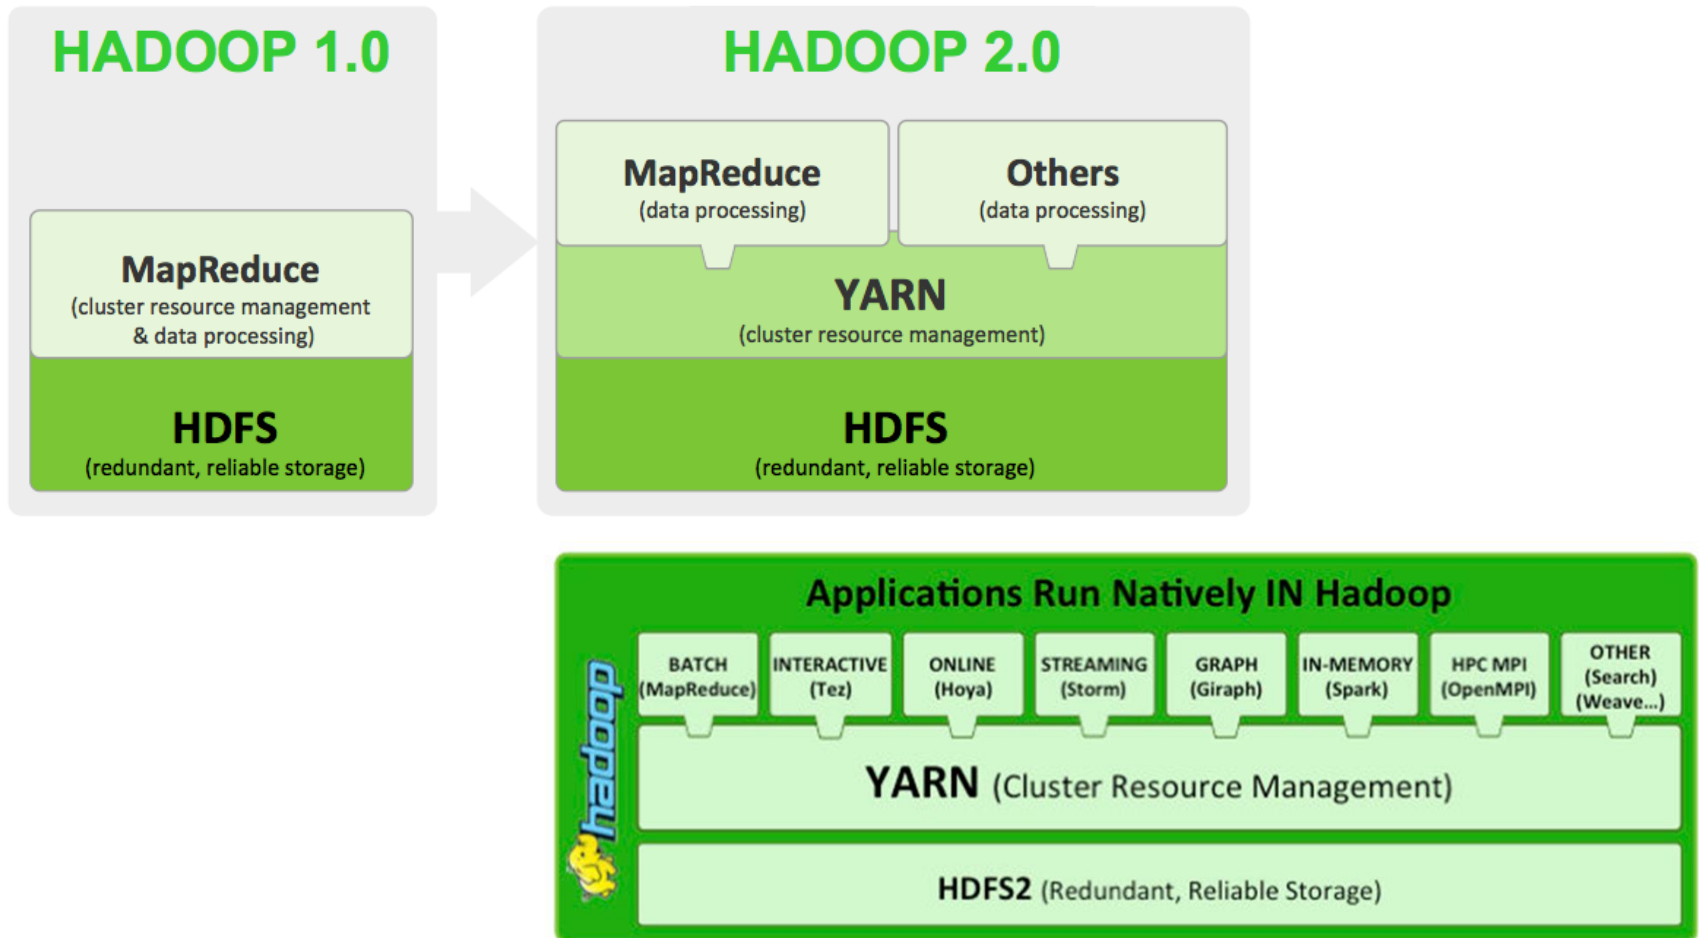
- **Job Tracker**

- Schedules and manages jobs

- **Task Tracker**

- Executes MapReduce jobs on individual data node

Hadoop 2.0



Setup

- **Three options:**
 - Standalone (single Java process)
 - **Pseudo-Distributed (separate Java processes)**
 - Fully-Distributed
- **Prerequisites (on virtual machine):**
 - Ubuntu server 12.04
 - Ubuntu desktop (for monitoring)
 - Oracle (Sun) Java 1.7.0_25
 - <http://www.webupd8.org/2012/01/install-oracle-java-jdk-7-in-ubuntu-via.html>
 - SSH

Setup Files

- **hosts (ip address)**
- **.bashrc (Java dir, home dir)**
- **hadoop configuration files in /usr/local/hadoop/conf**
 - hadoop-env.sh (Java dir)
 - hdfs-site.xml (replication factor)
 - mapred-site.xml (host/port for jobtracker)
 - master/slave (ips for multi-node cluster)

Login

- **User: rmacc**
- **Password: rmacc**
- **Start terminal (Ctrl+Alt+F1)**
- **Login as hduser**
 - User: **hduser**
 - Password: **hduser**
- **Change directory to hadoop**
 - with **cd \$HADOOP_PREFIX**
 - or **cd /usr/local/hadoop**



Starting Hadoop Daemons

- **All:**

- \$ bin/start-all.sh

- Log output in *logs* directory

Check Daemons

- **jps** →
 - 1367 Jps
 - 8695 DataNode
 - 8609 NameNode
 - 6318 SecondaryNameNode
 - 2600 JobTracker
 - 2830 TaskTracker


Web Interfaces

- <http://localhost:50030> → Cluster status and jobs
- <http://localhost:50070> → HDFS

HDFS

- **Hadoop Distributed File System**

Ubuntu (Linux)



Hadoop



```
bin/hadoop dfs -copyFromLocal mis6110/Files/file.txt mis6110/KDD/file.txt
```

```
bin/hadoop dfs -copyToLocal mis6110/Movies/file.txt mis6110/Files/file.txt
```

local file system
/usr/local/hadoop/...

mis6110

Files

Scripts

(HDFS)
/user/hduser/...

mis6110

KDD

Movies

Create
Directory

mkdir

Show content
of directory

ls

bin/hadoop dfs -mkdir

bin/hadoop dfs -ls

Create
Directory

sh mkdir

Show content
of directory

sh ls

fs -mkdir

fs -ls



HDFS Shell

- **bin/hadoop dfsadmin -help → all admin commands**
- **bin/hadoop dfs -help → all commands**
- **Most commands similar to unix**
 - dfs -copyFromLocal
 - dfs -ls
- Shell commands:
<http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

Importing Data

- **KDD Example:** Legcare sales data from [KDD Cup 2000](#)

“We wish to thank Blue Martini Software for contributing the KDD Cup 2000 data”

- Cleaned for easier/faster use
- Copy file KDDCupCleaned.txt to hdfs

Copying File to HDFS

- **Create new directory:**

```
$ bin/hadoop dfs -mkdir rmacc
```

- **Copy files:**

```
$ bin/hadoop dfs -copyFromLocal ../rmacc/Beowulf.txt  
rmacc/
```

```
$ bin/hadoop dfs -copyFromLocal ../rmacc/KDD*  
rmacc/
```

- **Check:**

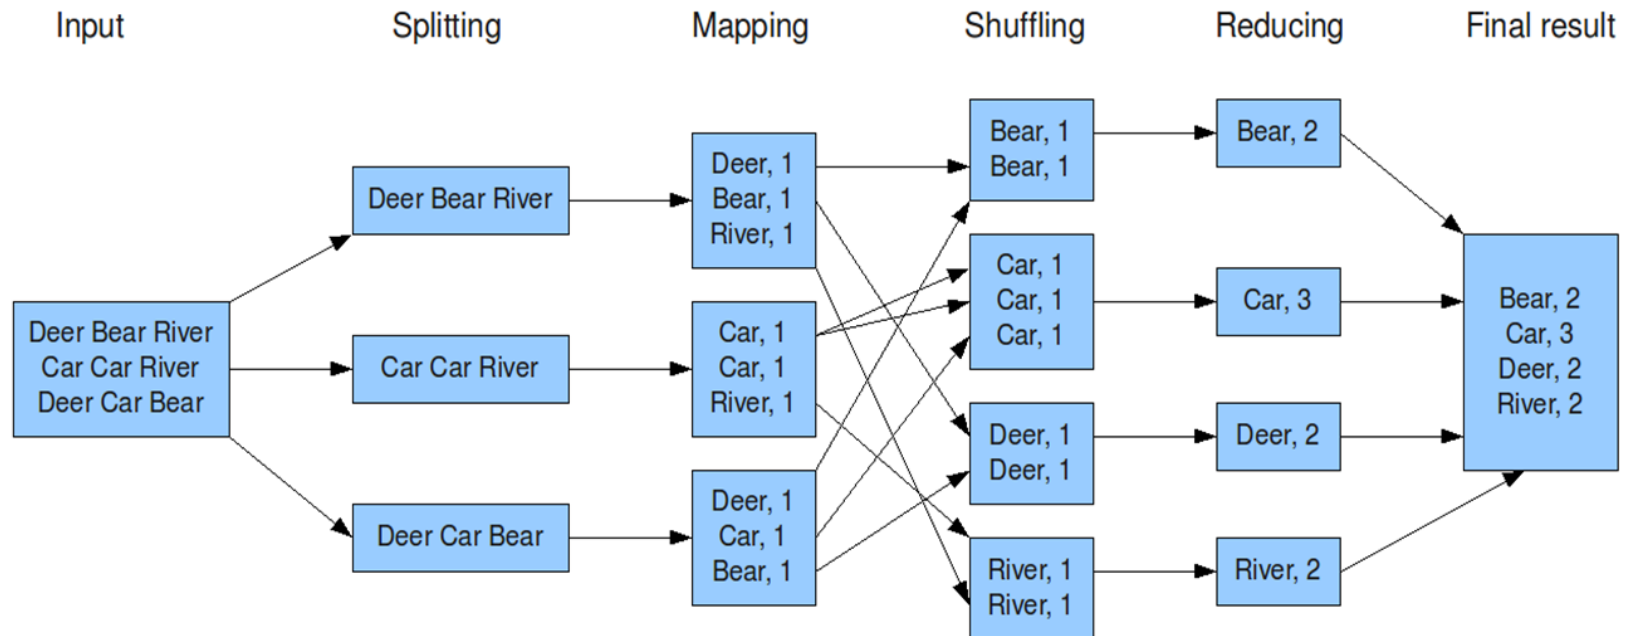
```
$ bin/hadoop dfs -lsr rmacc
```

- ***Localhost:50070 → Browse the filesystem***



MapReduce WordCount Example

The overall MapReduce word count process



Java Code for WordCount Example

```
1. package org.myorg;
2.
3. import java.io.IOException;
4. import java.util.*;
5.
6. import org.apache.hadoop.fs.Path;
7. import org.apache.hadoop.conf.*;
8. import org.apache.hadoop.io.*;
9. import org.apache.hadoop.mapred.*;
10. import org.apache.hadoop.util.*;
11.
12. public class WordCount {
13.
14.     public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
15.         private final static IntWritable one = new IntWritable(1);
16.         private Text word = new Text();
17.
18.         public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
19.             String line = value.toString();
20.             StringTokenizer tokenizer = new StringTokenizer(line);
21.             while (tokenizer.hasMoreTokens()) {
22.                 word.set(tokenizer.nextToken());
23.                 output.collect(word, one);
24.             }
25.         }
26.     }
27.
28.     public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
```

Java Code for WordCount Example


```
29.     public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
30.         int sum = 0;
31.         while (values.hasNext()) {
32.             sum += values.next().get();
33.         }
34.         output.collect(key, new IntWritable(sum));
35.     }
36. }
37.
38. public static void main(String[] args) throws Exception {
39.     JobConf conf = new JobConf(WordCount.class);
40.     conf.setJobName("wordcount");
41.
42.     conf.setOutputKeyClass(Text.class);
43.     conf.setOutputValueClass(IntWritable.class);
44.
45.     conf.setMapperClass(Map.class);
46.     conf.setCombinerClass(Reduce.class);
47.     conf.setReducerClass(Reduce.class);
48.
49.     conf.setInputFormat(TextInputFormat.class);
50.     conf.setOutputFormat(TextOutputFormat.class);
51.
52.     FileInputFormat.setInputPaths(conf, new Path(args[0]));
53.     FileOutputFormat.setOutputPath(conf, new Path(args[1]));
54.
55.     JobClient.runJob(conf);
57. }
58.
```





There Must be an Easier Way

Pig (Latin)



Pig

- **High-level data processing language (Pig Latin)**
- **Resides on user machine, not cluster**
- **Pig Latin compiled into efficient MapReduce jobs**

Test Pig Installation

- **(Hadoop has to be running)**
 - \$ cd \$PIG_HOME
 - \$ bin/pig -help

How to Run Pig

- **Grunt interactive shell**
 - Two modes:
 - *Local, standalone (pig -x local)*
 - *Hadoop, distributed (pig -x mapreduce)*
- **Scripts (.pig)**
- **Embedded in Java or Python**
- **PigPen, Eclipse plugin**

Pig Statements in Grunt

- **LOAD** → **Transform data** → **DUMP** or **STORE**
 - **Example:**
 - *grunt> A = load 'student' using PigStorage() AS (name:chararray, age:int, gpa:float);*
 - *grunt> B = foreach A generate name;*
 - *grunt> dump B*
- “A “ is called a “relation” or “outer bag”

Pig Example: KDD Data

Key	Date	Time	Unit Price	Order LineID	Qty	Order Status	Tax	Amount	Weekday	Hour	City	State	Customer ID
1	2/27/2000	08:06:35	15	1	15	Shipped	1.3	16.3	Sunday	8	Westport	CT	62
2	3/30/2000	10:00:18	9	1	9	Shipped	0	9	Thursday	10	Westport	CT	62
3	1/28/2000	14:43:34	12	1	12	Shipped	1.02	13.02	Friday	14	San Francisco	CA	96
4	1/29/2000	10:22:37	12	1	12	Shipped	0.87	12.87	Saturday	10	Novato	CA	132
5	2/1/2000	08:44:48	6.5	1	6.5	Shipped	0.55	7.05	Tuesday	8	Cupertino	CA	168
6	2/29/2000	10:31:42	15	1	15	Shipped	1.24	16.24	Tuesday	10	San Ramon	CA	184
7	2/29/2000	10:31:42	14	1	14	Shipped	1.16	15.16	Tuesday	10	San Ramon	CA	184
8	2/29/2000	10:31:42	6.5	2	6.5	Shipped	1.07	14.07	Tuesday	10	San Ramon	CA	184
9	3/8/2000	16:48:47	11	3	11	Shipped	2.72	35.72	Wednesday	16	San Ramon	CA	184
10	1/30/2000	14:13:57	10	1	10	Shipped	0	10	Sunday	14	Scarsdale	NY	224
11	1/30/2000	14:13:57	13.5	1	13.5	Shipped	0	13.5	Sunday	14	Scarsdale	NY	224
12	2/26/2000	03:42:17	12.7	1	12.7	Shipped	0	12.7	Saturday	3	Novato	CA	236
13	3/30/2000	11:51:44	6.5	1	4.88	Shipped	0	4.88	Thursday	11	Novato	CA	236
14	3/30/2000	11:51:44	6.5	1	4.88	Shipped	0	4.88	Thursday	11	Novato	CA	236
15	3/30/2000	11:51:44	7	1	7	Shipped	0	7	Thursday	11	Novato	CA	236
16	3/30/2000	11:51:44	12	1	12	Shipped	0	12	Thursday	11	Novato	CA	236

Pig LOAD Function

- **“Pigs eat anything”**
- **LOAD 'data' [USING function] [AS schema];**
- **USING**
 - PigStorage → structured text file (default)
 - TextLoader → unstructured UTF-8 data
 - Other and User Defined Functions
- **AS**
 - (Field1[:type], Field2[:type], ... FieldX[:type])
 - Bytearray default type

Pig Example: Loading Data

> pig -x mapreduce

> a = LOAD 'rmacc/KDDCupCleaned.txt';

> All Statements end with semicolon !!!

Pig Debugging Statements

Debug Operator	Description
DUMP	Display results
DESCRIBE	Display schema of relation
EXPLAIN	Display execution plan
ILLUSTRATE	Display step-by-step execution
<i>Full list</i>	http://pig.apache.org/docs/r0.11.1/test.html#diagnostic-ops

- **Describe and illustrate only work if schema is provided**

Pig Example: Loading Data

```
> a = LOAD 'rmacc/KDDrmacc' , AS  
      (key,date,time, qty:float);
```

- **Columns can be accessed by**
 - \$0 for second column (first contains key)
 - or name (key, date, time....)

Pig Data Types

- **Simple:**

- int, long, float, double, chararray, bytearray, boolean

- **Complex:**

- tuple - a set of fields (10, 5, alpha)
- bag - a collection of tuples {(10,5,alpha) (8,2,beta)}
- map - a set of key value pairs [key#value]

Pig Example: Reduce # of Fields

- > a = LOAD 'rmacc/KDDCupCleaned.txt' AS
(key:int,date,time,up,ol,qty,os,tax,amount:float,
wd,hour,city,state,ci);
- > b = FOREACH a GENERATE key, date, time,
amount;
- > STORE b INTO 'rmacc/KDDCupShort'.txt
USING PigStorage(',') ;

Pig Example: Group per Date

- > a = LOAD 'rmacc/KDDrmacc.txt' AS
(key,date,time,qty:float);
- > groupday = GROUP a BY date;
- > illustrate groupday →
 - > *group* is new key for each bag (day)
 - > tuples within data within each bag

Pig Example: Sum per Date

- > a = LOAD 'rmacc/KDDrmacc.txt' AS (key,date,time,qty:float);
- > groupday = GROUP a BY date;
- > sumday = FOREACH groupday GENERATE group, SUM(a.qty); sr
- > STORE sumday INTO 'rmacc/sumday';

Evaluation Functions

(Case Sensitive)

Function	Description
AVG	Calculates average
CONCAT	Concatenates two expressions of identical type
COUNT	Counts the number of elements in a bag
COUNT_STAR	Like count by includes NULL values in count
DIFF	Compares two fields in a tuple
IsEmpty	Checks if a bag or map is empty
MAX	Calculates maximum
MIN	Calculates minimum
SIZE	Computes the number of elements (characters)
SUM	Calculates sum
TOKENIZE	Splits a string and outputs a bag of words
<i>List with examples</i>	http://pig.apache.org/docs/r0.11.1/func.html#eval-functions



Other Functions

- **Math Functions**
- **String Functions**
- **Datetime Functions**
- **Tuple, Bag, Map Functions**
- **User Defined Functions**

Pig Example: Filter Purchases > \$100

```
> a = LOAD 'rmacc/sumday' AS (date,sum:float);  
> bigpur = FILTER a BY sum>1000;  
> STORE bigpur INTO 'rmacc/bigpur';
```

Relational Operators

Operators	Description
LOAD	Loads data from the file system
GROUP	Groups the data in one or more relations
FOREACH	Generates data transformations based on columns of data
FILTER	Selects tuples from a relation based on some condition
<i>Full list with examples</i>	http://pig.apache.org/docs/r0.11.1/basic.html#Relational+Operators



Other Operators

- **Arithmetic Operators**
- **Boolean Operators**
- **Cast Operators**
- **Comparison Operators**
- **Type Construction Operators**
- **Dereference Operators**
- **Disambiguate Operator**
- **Flatten Operator**
- **Null Operators**
- **Sign Operators**

Pig WordCount Code

- > b = LOAD 'rmacc/Beowulf.txt' AS be;
- > beowords = FOREACH b GENERATE
flatten(TOKENIZE(beo)) as bw;
- > wg = GROUP beowords BY bw;
- > wc = FOREACH wg GENERATE group,
COUNT(beowords) as bc;
- > sumord =ORDER wc BY bc;
- > STORE sumord INTO 'rmacc/beowulfwc'

Hadoop Summary

- **Accessible – runs on commodity hardware**
- **Robust – handles hardware failures**
- **Scalable – by adding more nodes**
- **Simple – allows users to quickly write efficient parallel code**
- **Pig - easy to learn**

Conclusions

- **Individual components “easy” to setup → integration more complicated**
- **Resources**
 - Apache Hadoop (download and docu)
 - <http://hadoop.apache.org>
 - Online Searches
 - Books for overview, not technical details
- **“Evolving project” → constantly changing, documentation can’t keep up with development**

Questions ?